Housing Price Prediction Project

● Introduction

➔ This project focuses on predicting housing prices using machine learning algorithms. The goal is to develop a model that accurately predicts the difference between population and profit per city by making an graph for population vs profit per city.

● Installation

➔ To run this project, you will need Python 3.x installed on your machine. Additionally, ensure you have the following libraries installed:

1. Pandas
2. NumPy
3. Scikit-learn
4. Matplotlib
5. Seaborn

● You can install these libraries using pip:

➔ pip install pandas numpy scikit-learn matplotlib seaborn

● Usage

➔Clone this repository to your local machine.

➔Navigate to the project directory.

➔Run the C1_W2_Linear_Regression.ipynb script using Python:

- ## Data

➔ Two dataset used in this project is included in the repository as ex1data1.txt and ex1data2. It contains information about population and city to find answer for population vs profit per city.

- ## Model

➔ The model used for predicting housing prices is Linear regression algorithm trained on the provided dataset. The model has been evaluated for accuracy and performance using appropriate metrics.

- ## Results

➔ Upon running the script, you will see the predicted prices of houses based on the input features. Additionally, visualizations may be generated to provide insights into the data and model performance.

- ## Contributing

➔ Contributions to this project are welcome. If you have suggestions for improvement or would like to add new features, please feel free to open an issue or submit a pull request.